



POLITECNICO
MILANO 1863

SCUOLA DI INGEGNERIA INDUSTRIALE
E DELL'INFORMAZIONE

ARTIFICIAL NEURAL NETWORKS AND DEEP LEARNING

Homework 2: Time Series Forecasting

Authors:

- DOMENICO CACACE
- MIGUEL GONZALEZ
- USEVALAD MILASHEUSKI

Team name: DMU

1. Introduction

Our strategy for this homework is to start with simple Recurrent Neural Networks, evaluating pros and cons of different architectures. Afterwards we implement a seq2seq model, on which we can then test *novel* methodologies, such as the Luong attention mechanism.

2. Dataset

The dataset used in this project consists of 7 time series, each containing 68529 samples. As for pre-processing, we normalized the data (MinMax normalization) and extracted the sequences to be fed to the network. To extract the sequences we employed a sliding window approach; for our models, we found that a window size of 200 and a stride of 5 yield the best results.

The models submitted to the competition are trained on the whole dataset, with a 80/20 split between training and validation; to determine the performance of a model (based on how well they predict the future) and decide what models to submit we reserved ~15% of the original dataset for testing.¹

We evaluated both oneshot and autoregressive predictions on most of our model, and found that the latter produced better predictions in almost all the cases.

¹This step became the only method to evaluate our models, due to the Codalab issues

3. Recurrent neural networks

As a first step, we evaluated the different recurrent layers that Keras offers. For the sake of consistency, we used the same model structure, consisting of:

- Recurrent layer(s)
- Dense layer
- Reshape layer
- Conv1D layer

For the recurrent layers we used LSTM and GRU; we did not consider at all the SimpleRNN layer, due to its vanishing gradient issue. We trained our models using the Mean Absolute Error and the Mean Squared Error as accuracy and loss functions respectively; we also employed early stopping (max 200 epochs) and learning rate reduction.

Tests with *unidirectional* layers yielded similar results for both LSTM and GRU layers; changing the number of units per layer (50, 100, 256) and the number of layers (2, 3, 4) did not bring a significant improvements. During this phase we observed that, in some cases, there was a jump discontinuity between the last input and the first predicted value: to solve this issue we employed a dropout layer after every recurrent one, and observed that the gap disappeared.




Figure 1: no dropout




Figure 2: dropout (0.2)

3.1. Bidirectional layers

With bidirectional layers we observed significant differences among the models tested; as above, we changed the number of units per layer and the number of recurrent layers. We observed that the BiGRU models had significantly worse performances than BiLSTM ones for every variation tested: we attributed these outcomes to the structure of the memory units, as GRUs have less parameters than LSTM units.

Regarding BiLSTM models, we obtained the best results with 3 recurrent layers, while increasing the number of layers yields slightly worse results.




Figure 3: 3 BiLSTM layers (256 units), 0.2 Dropout

4. To seq2seq and beyond

Disclaimer: all the models presented in this section have been developed after the CodaLab outage; the results we present have been obtained via cross validation, not on the secret test set.

Our next step consists in building an encoder/decoder model. Instead of starting from scratch, we employed the results obtained in the previous experiments; following the classic encoder/decoder structure, we have:

- 3 BiLSTM layers, for the encoder
- 1 `RepeatVector`, to connect the components
- 2 BiLSTM layers, for the decoder
- 1 `TimeDistributed Dense` output layer

From our tests we can see results slightly better, on average, to the ones obtained with the BiLSTM model; this did not come as a surprise to us, but gave us a base for the final step.

4.1. Luong attention

As a final step, we implemented the Luong attention mechanism; in this phase we had some trouble with the `Attention` layer provided by Keras, so we implemented it manually with dot products and Softmax (see the code and the model scheme in the appendix for more details).

The results obtained with this last model are, on average, better than the previous: we can note how the gap between the predicted and actual value shrinks, e.g. in proximity of the minima of *Hype Root* and *Crunchiness*, or stays almost the same.




Figure 4: Encoder/Decoder with Luong attention

5. Conclusions

In this report we presented our small journey in the world of time series forecasting with neural network architectures. Starting from the basic memory blocks we built and evaluated several models, spacing between different architectures. In the end we found that the encoder/decoder architecture, in particular with the Luong attention mechanism, yielded the best results in our scenario.

Appendix A: *BiLSTM* model

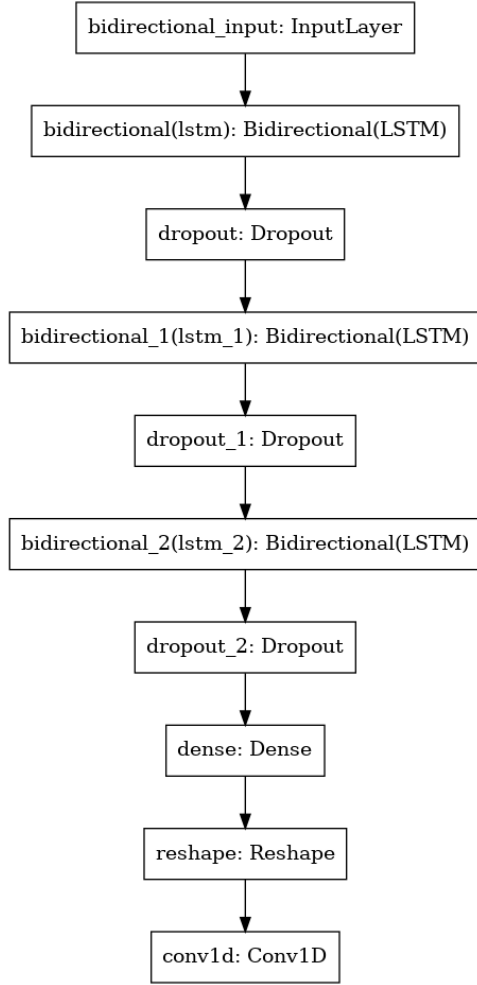


Figure 5: Encoder/Decoder with Luong attention

Appendix B: *attention* model

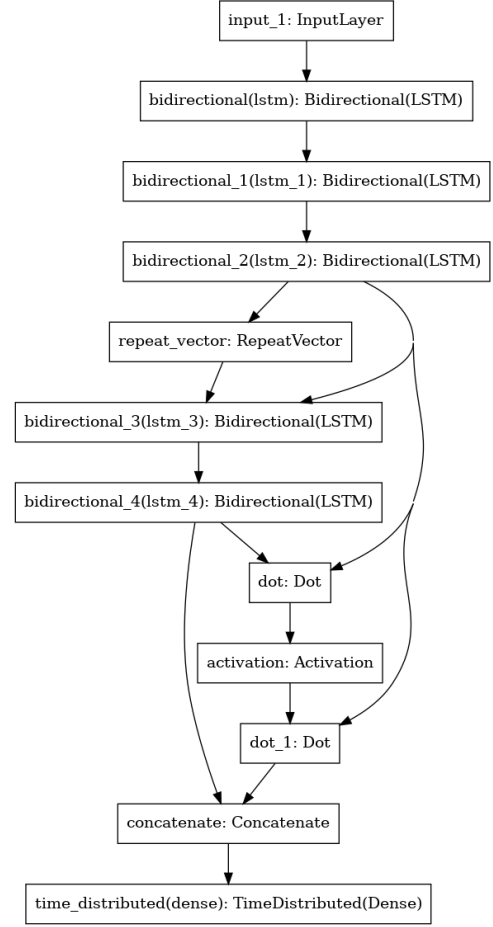


Figure 6: Encoder/Decoder with Luong attention